

Zahid Ahmed

SOFTWARE DEVELOPER · AI EVALUATION & BENCHMARK DESIGN

Malegaon, Maharashtra, India | +91 86690 82139 | zahid23saim@gmail.com
[linkedin.com/in/zahid-ahmed-5a4129424](https://www.linkedin.com/in/zahid-ahmed-5a4129424) | github.com/zahid23saim | zahid23saim.github.io | dev.to/zahid23saim

PROFESSIONAL SUMMARY

Software developer and AI-evaluation specialist with an MCA, 3+ years of Python development, and 15+ years in data quality and research. As a Handshake AI Fellow on Project Dynamo, builds Terminal-Bench-style tasks — Docker environments, reference solutions, and deterministic pytest verifiers — that pinpoint where frontier AI coding agents fail. Brings full-stack experience in Python, .NET/C#, SQL, and Oracle PL/SQL, plus multilingual LLM and speech-output evaluation (English, Hindi, Tamil, Marathi) from Appen.

CORE SKILLS

AI Evaluation	Agentic benchmark task design, pytest verifiers, pass@k calibration, failure-trajectory analysis, LLM output evaluation, factuality & STEM auditing, TTS quality (MOS), data annotation
Programming	Python (3+ yrs), C#, Java, SQL, PL/SQL, JavaScript, Bash, C, C++
Frameworks & Web	.NET / .NET Core, ASP.NET Core, Entity Framework, MVC, HTML, CSS
Tools & Platforms	Docker, Git & GitHub (PRs, CI), Linux CLI, uv, Harbor eval harness, pytest, xUnit, MS Excel
Databases	Oracle (DBA & PL/SQL), MySQL, MariaDB, stored procedures, query optimization
Practices	OOP & system design, automated testing, debugging, technical writing
Spoken Languages	English, Hindi, Tamil, Marathi — professional fluency

PROFESSIONAL EXPERIENCE

AI Fellow, Project Dynamo — Handshake AI	Jul 2026 – Present
- Build end-to-end, Terminal-Bench-style tasks that test frontier AI coding agents in real terminal environments: a pinned Docker image, a reference (oracle) solution, and a deterministic pytest verifier. - Design expert-level, multi-step problems drawn from realistic professional work, novel enough to resist memorization and graded only on the artifacts an agent produces, not the path it takes. - Write tight task specifications ($\leq 1,500$ tokens) whose numbered success criteria map one-to-one to verifier tests; harden verifiers against hard-coding, shortcuts, and reward hacking. - Calibrate tasks in the Harbor harness (oracle passes, no-op agent fails), then prove difficulty through staged pass@k trials — a task qualifies only if a frontier agent fails at least half its runs for genuine reasoning errors. - Ship through a GitHub fork-and-PR workflow gated by automated CI (schema validation, agent trials, LLM rubric review) and two rounds of expert human review.	
AI Data Contributor & Evaluator — Appen	May 2025 – Jun 2026
- Evaluated LLM-generated responses across multilingual datasets (English, Hindi, Tamil, Marathi), writing detailed feedback to sharpen linguistic nuance and technical precision. - Ran Python-assisted factuality and quality checks on STEM-focused AI queries and classified educational content. - Rated AI audio dubbing and text-to-speech (TTS) output on Mean Opinion Score (MOS) scales for naturalness, prosody, and translation accuracy.	
Freelance Translator & Transcriptionist — Self-employed	May 2024 – Apr 2025
- Delivered verbatim transcription and technical/general translation across English, Hindi, Tamil, and Marathi under strict notation guidelines. - Audited audio against transcripts to keep every deliverable accurate, consistent, and client-ready.	

Software Developer — Emsys Solutions Pvt Ltd

2021 – 2023

- Designed, built, and maintained custom software for vehicle showrooms and dealerships, streamlining sales and day-to-day operational workflows.
- Developed web applications and backend services with Python, .NET Core, C#, Entity Framework, MVC, and JavaScript.
- Wrote and optimized complex SQL, MySQL, and MariaDB stored procedures for fast, reliable data access.
- Applied OOP principles to improve application architecture, scalability, and maintainability.

Datamatics Business Solutions Ltd

Apr 2009 – 2024

Quality Analyst

Jul 2018 – 2024

- Analyzed client requirements to produce accurate effort estimates, resource plans, and turnaround-time (TAT) commitments.
- Maintained project logs and resource-utilization reports to balance workload and improve team efficiency.

Senior Research Associate

Apr 2016 – Jun 2018

- Led secondary-research initiatives, validating entity data with advanced web-search techniques.
- Coordinated workflow and task allocation to meet client deadlines consistently.

Research Associate

Apr 2009 – Apr 2015

- Ran quality-control checks on research datasets and trained new hires on process and Excel data formatting.

PROJECTS & OPEN SOURCE

Block & Bell JavaScript · Cloudflare Workers · Durable Objects · WebSockets · [live demo](#)

Real-time railway-signaling game for 2–4 players where each player's orders appear only on a neighbor's screen. A server-authoritative Durable Objects backend sends each seat only its own view; a seeded generator solves every shift before printing it (verified on 600 generated layouts), and a capacity fix raised simulated completion from 2/400 to 398/400 nights.

llm-eval-harness Python · pytest · [live demo](#)

Dependency-free harness that scores LLM answers against a gold set, with per-question match rules (exact, contains, numeric), gold-file validation, actionable failure output, and a CI-friendly exit code.

aspnet-minimal-api C# · ASP.NET Core (.NET 8) · xUnit

Single-file task-tracker Web API with full CRUD, input validation, and correct HTTP status codes, covered by in-memory integration tests (WebApplicationFactory) that exercise the real endpoints — 7/7 passing.

plsql-bulk-examples Oracle PL/SQL

Runnable examples that turn slow row-by-row loads into batched BULK COLLECT / FORALL processing with SAVE EXCEPTIONS error handling — an order-of-magnitude speedup on a 200k-row table.

Smart Maintenance Portal Java · Python · MySQL · HTML/CSS

Full-stack municipal-complaint portal with an admin dashboard that routes complaints to departments and tracks resolution status in real time.

Technical Writing dev.to · 5 articles

Articles on catching LLM regressions with a CI gate, automating LLM answer evaluation in Python, speeding up PL/SQL with BULK COLLECT and FORALL, building a tested ASP.NET Core Minimal API, and evaluating LLMs across four languages.

EDUCATION & CERTIFICATIONS

Master of Computer Applications (MCA) — Savitribai Phule Pune University

2009 – 2012

B.Sc. in Information Technology — Akbar Peerbhoy College

2005

Oracle DBA & PL/SQL Certification — Karrox Technology

2007